



picoJava™: **A Hardware Implementation** **of the Java Virtual Machine**

Marc Tremblay and Michael O'Connor
Sun Microelectronics

The Java – picoJava Synergy

- **Java's origins lie in improving the consumer embedded market**
- **picoJava is a low cost microprocessor dedicated to executing Java™-based bytecodes**
 - Best system price/performance
- **It is a processor core for:**
 - Network computer
 - Internet chip for network appliances
 - Cellular phone & telco processors
 - Traditional embedded applications



Java in Embedded Devices

Products in the embedded market require:

- **Robust programs**
 - Graceful recovery vs. crash
- **Increasingly complex programs with multiple programmers**
 - Object-oriented language and development environment
- **Re-using code from one product generation to the next**
 - Portable code
- **Safe connectivity to applets**
 - For networked devices (PDA, pagers, cell phones)

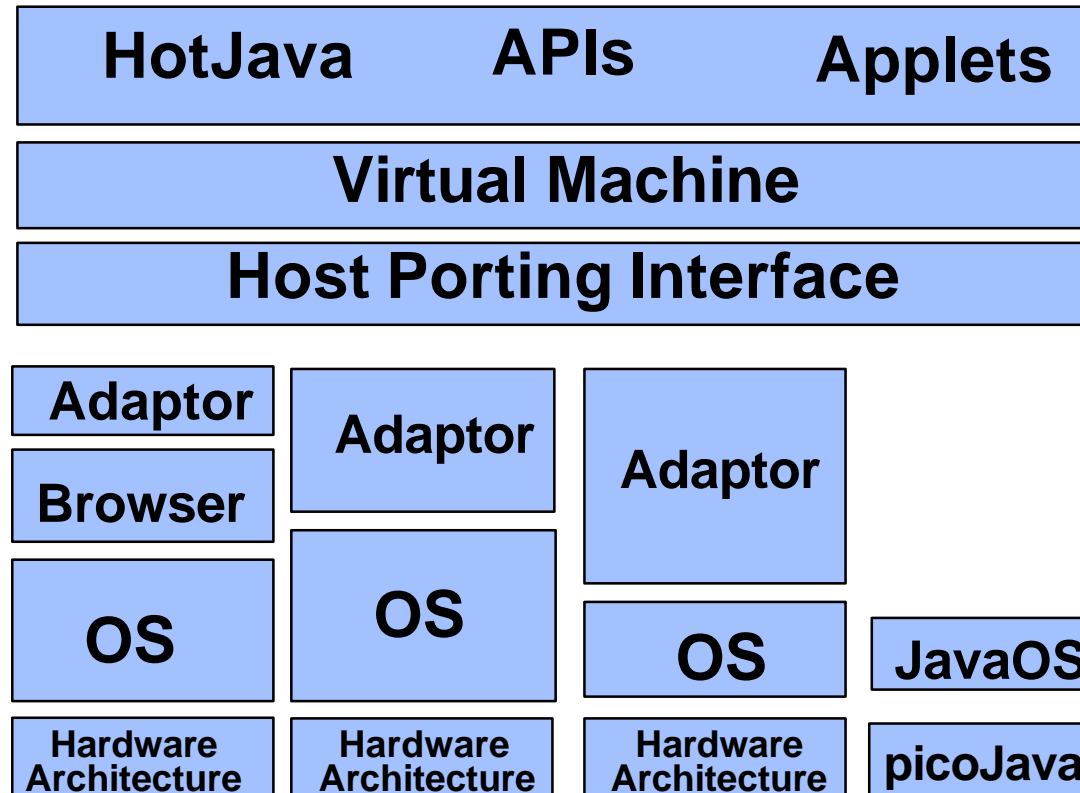


Important Factors to Consider in the Embedded World

- **Low system cost**
 - Processor, ROM, DRAM, etc.
- **Good performance**
- **Time-to-market**
- **Low power consumption**



Various Ways of Implementing the Java Virtual Machine



picoJava

- **Directly executes bytecodes**

- Excellent performance
- Eliminates the need for an interpreter or a JIT compiler
- Small memory footprint

- **Simple core**

- Legacy blocks and circuits are not present

- **Hardware support for the runtime**

- Addresses overall system performance



Java Virtual Machine

■ What the virtual machine specifies:

- Instruction set
- Data types
- Operand stack
- Constant pool
- Method area
- Heap for runtime data
- Format of the class file



Virtual Machine —Instruction Set

- **Data types:** byte, short, int, long float, double, char, object, returnAddress
- **All opcodes have 8 bits, but are followed by a variable number of operands (0, 1, 2, 3, ...)**
- **Opcodes**
 - 200 assigned
 - 25 quick variations
 - 3 reserved

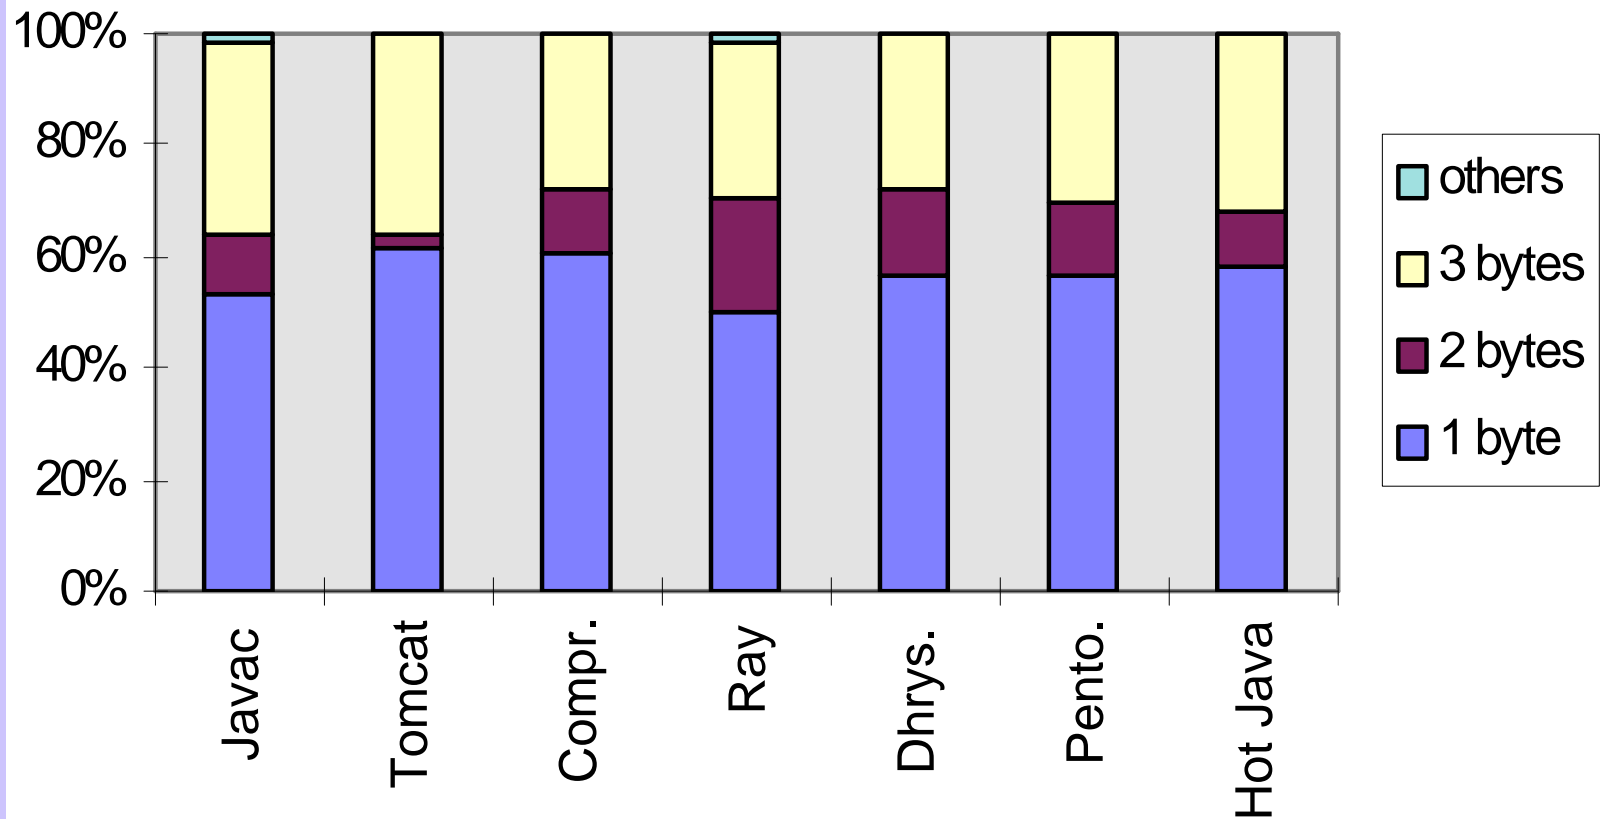


Java Virtual Machine Code Size

- **Java™-based bytecodes are small**
 - No register specifiers
 - Local variable accessed relative to a base pointer (VARS)
- **This results in very compact code**
 - Average JVM instruction is 1.8 bytes
 - RISC instructions typically require 4 bytes



Instruction Length



Java Virtual Machine Code Size

- Java bytecodes are about 2X smaller than the RISC code from the C++ compiler
- A large application (2500+lines) coded in both the C++ and Java languages



JVM – Instruction Set – RISCy

■ Some instructions are simple

bipush value	:push signed integer
iadd	:integer add
fadd	:single float add
ifeq	:branch if equal to 0
iload offset	:load integer from :local variable

JVM – Instruction Set – CISCy

■ Some instructions are complex

lookupswitch: “traditional” switch statement

byte 1	byte 2	byte 3	byte 4
opcode (171)	0..3 byte padding		
default offset			
numbers of pairs that follow (N)			
match 1			
jump offset 1			
match 2			
jump offset 2			
...			
...			
match N			
jump offset N			

Interpreter Loop

loop: 1: fetch bytecodes

2: indirect jump to

emulation code



Emulation Code

1: get operands

2: perform

operation

3: increment PC

4: go to loop



JVM: Stack-Based Architecture

- Operands typically accessed from the stack, put back on the stack
- Example — integer add:
 - Add top 2 entries in the stack and put the result on top of the stack
 - Typical emulation on a RISC processor

```
1: load tos
2: load tos-1
3: add
4: store tos-1
```

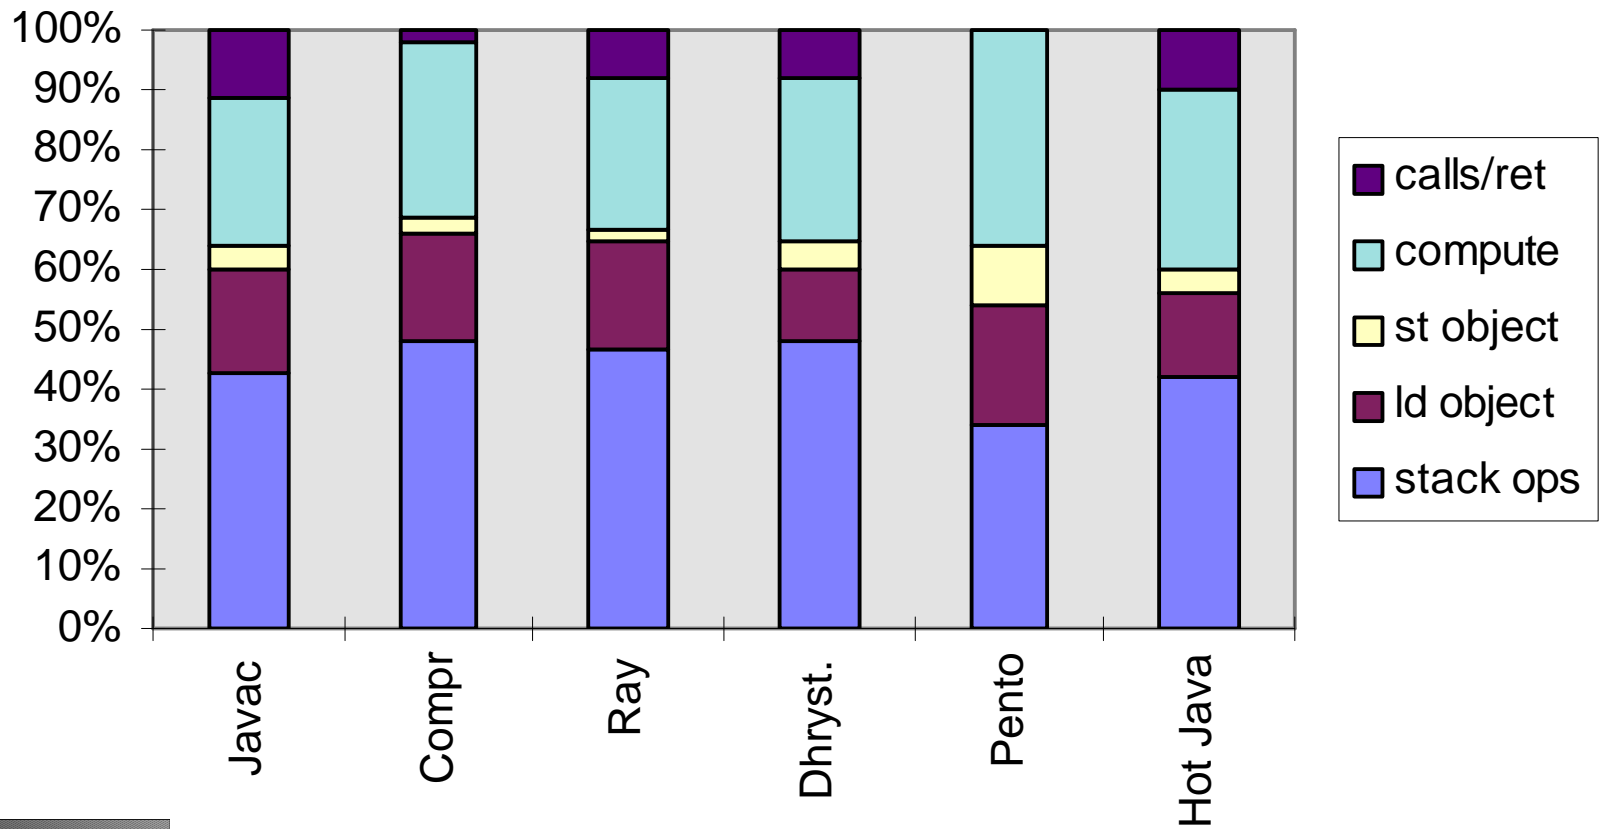


How to Best Execute Bytecodes?

- **Leverage RISC techniques developed over the past 15 years**
- **Implement in hardware only those instructions that make a difference**
 - Trap for costly instructions that do not occur often
 - State machines for high frequency/medium complexity instructions

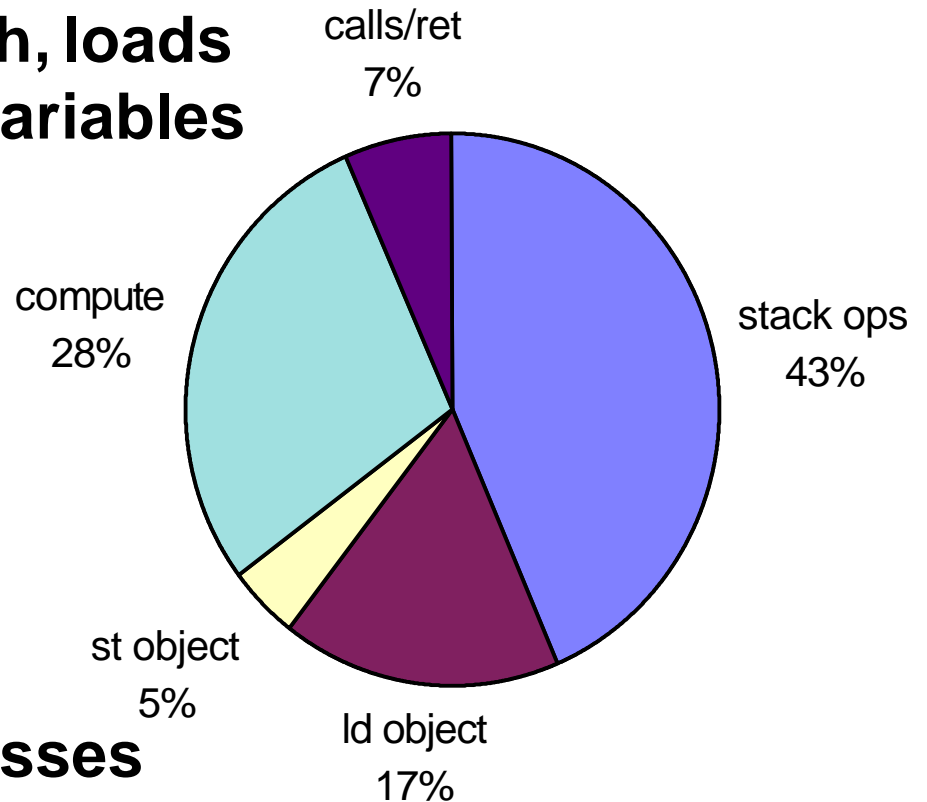


Dynamic Instruction Distribution

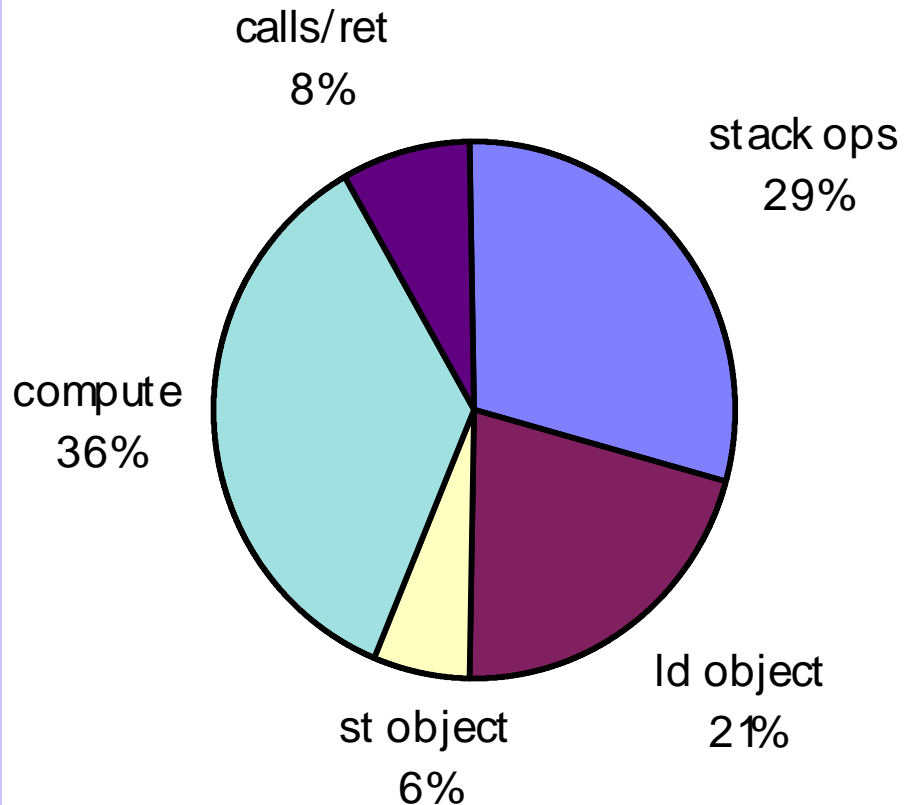


Composite Instruction Mix

- **Stack ops: dup, push, loads and stores to local variables**
- **compute: ALU, FP, compute branches**
- **calls/ret: method invocation virtual and non-virtual**
- **ld/st object: access to objects on the heap and array accesses**



Loads from Local Variables



- **Loads from local variables move data within the chip**
- **Target register is often consumed immediately**
- **Up to 60% of them can be hidden**
- **Resulting instruction distribution looks closer to a RISC processor**

Pipeline Design

■ RISC pipeline attributes

- Stages based on fundamental paths (e.g. cache access, ALU path, registers access)
- No operation on cache/memory data
- Hardwire all simple operations

■ Enhance classic pipeline

- Support for method invocations
- Support for hiding loads from local variables



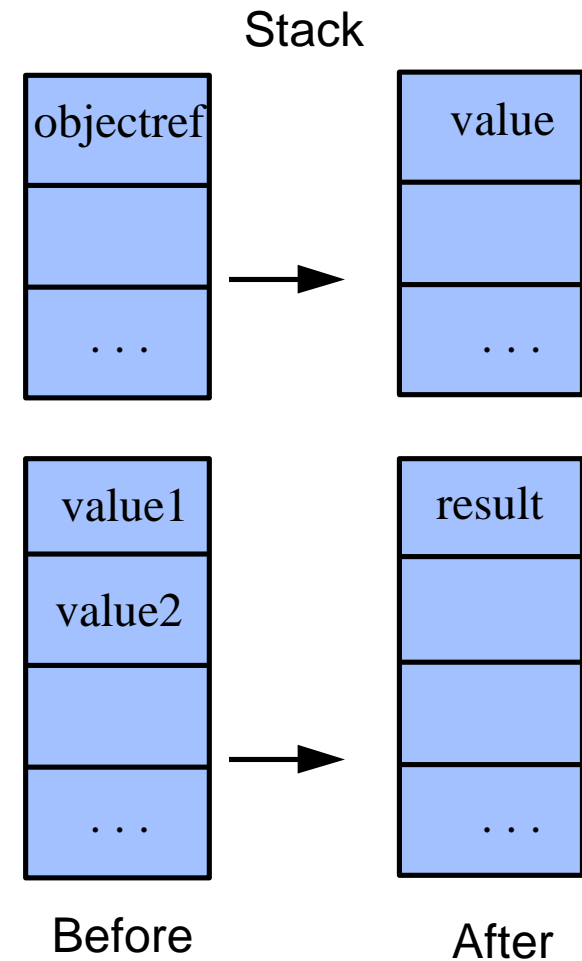
Implementation of Critical Instructions

`getfield_quick offset`

- Fetch field from object
- Executes as a “load [object + offset]” on picoJava

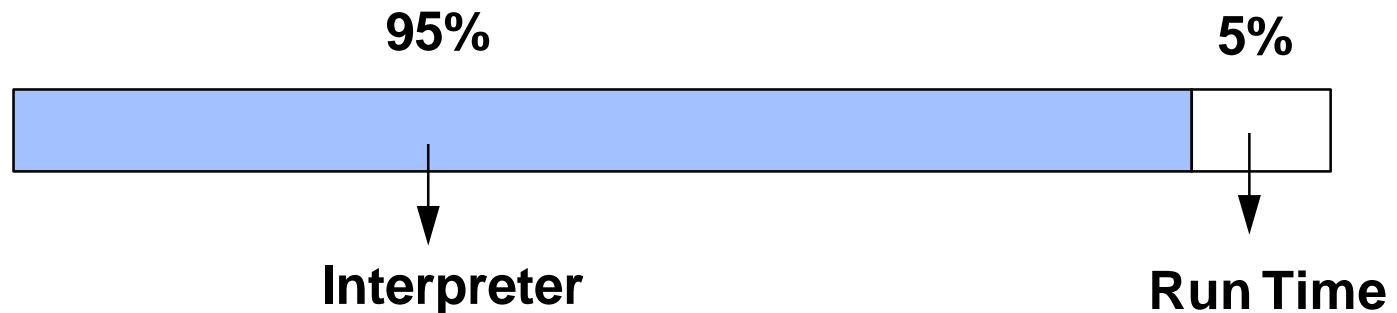
`iadd`

- Fully pipelined
- Executes in a single cycle



Typical Small Benchmarks (Caffeinemarks, Pentonimo, etc.)

■ Few objects, few calls, few threads



Speeding up the

Interpreter by 30X results in:

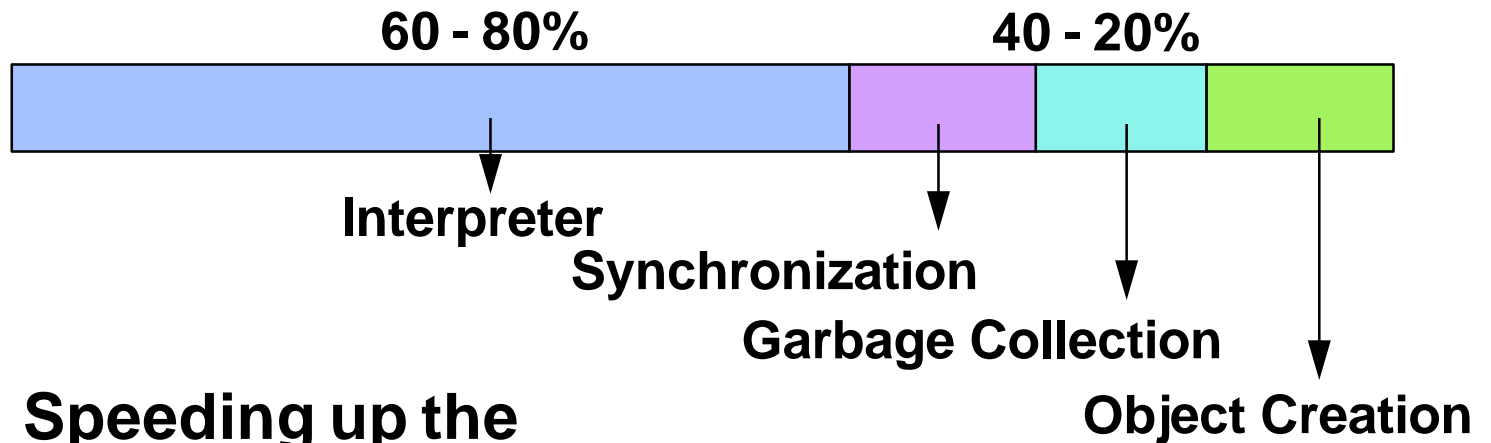
$$\begin{array}{rcl} 95 & \rightarrow & 3.2 \\ 5 & \rightarrow & 5 \\ \hline & & 8.2 \end{array}$$

=> Speedup of ~12X



Representative Applications

- Lots of Objects
- Threaded Code



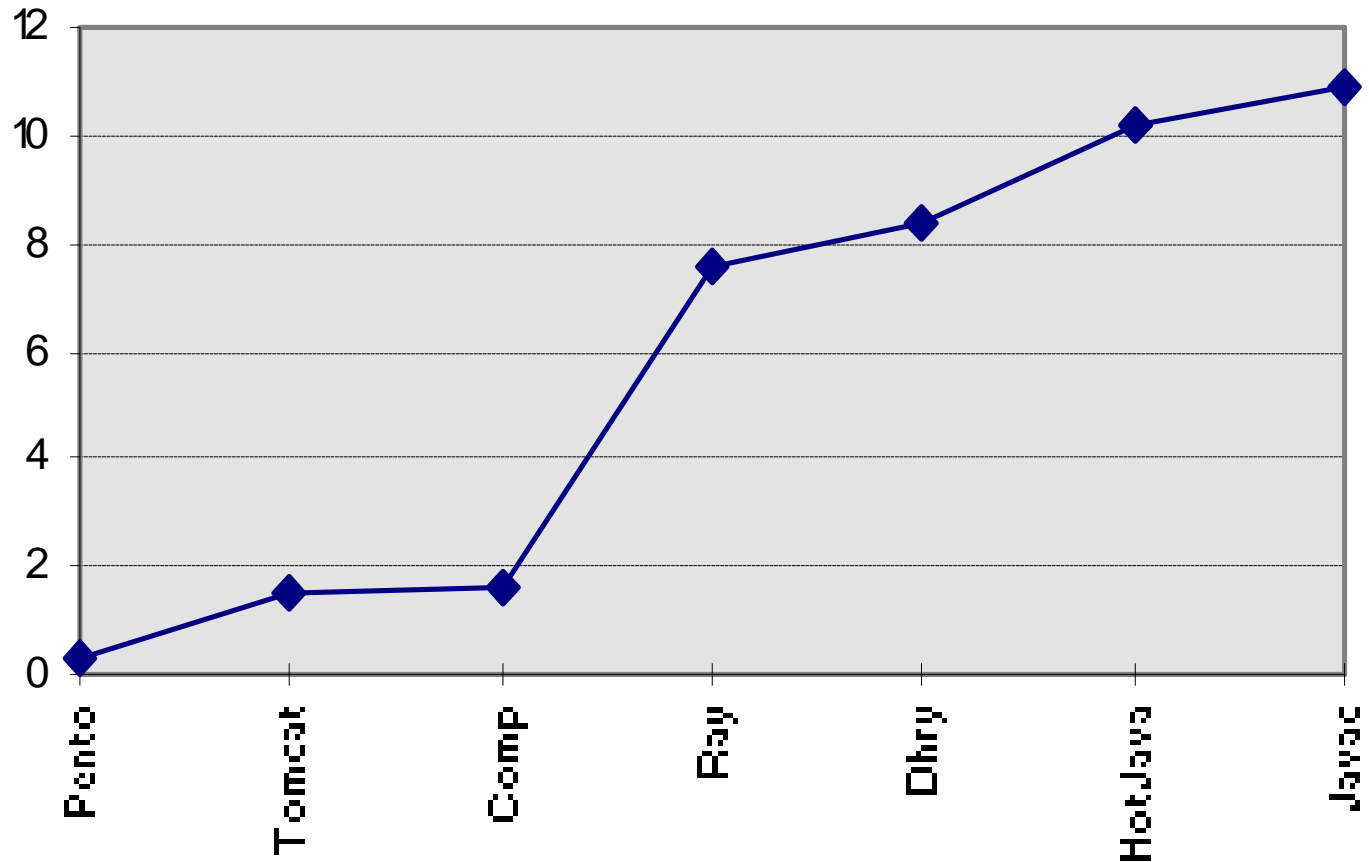
Speeding up the
Interpreter by 30X results in:

60	→	2
40	→	40
		<hr/>
		42

=> Speedup of ~2X



Percentage of Calls



Varies dramatically according to benchmark type

picoJava: A System Performance Approach

- **Accelerates object-oriented programs**

- simple pipeline with enhancements for features specific to bytecodes
- support for method invocation

- **Accelerates runtime**

(gc.c, monitor.c, threadruntime.c, etc.)

- Support for threads
- Support for garbage collection

- **Simple but efficient, non-invasive, hardware support**



System Programming

- Instructions added to support system programming
 - available only “under the hood”
 - operating system functions
 - access to I/O devices
 - access to the internals of picoJava



picoJava - Summary

Best system price/performance for running Java™-powered applications in embedded markets

- Embedded market very sensitive to system cost and power consumption
- Interpreter and/or JIT compiler eliminated
- Excellent *system* performance
- Efficient implementation through use of the same methodology, process and circuit techniques developed for RISC processors

